

WHAT WE BUILT

Type a topic. Get a teaching video.

Someone types one line — *"how do I handle an angry parent on the phone"* — and the system researches it, writes a twenty-beat script, reviews that script against our own standard, redrafts it until it passes, then draws it, narrates it, renders it, scores it, and stops to let a person watch it before anything is published. **One video at a time, and a human decides whether it ships.**

Why one video, not a course

The earlier version of this document described a course builder: type a subject, get nine lessons. That is not what exists, and the reason is worth stating plainly.

Three real course builds were attempted. They cost money, took hours, and produced no finished lesson — every one died somewhere in the middle, and because nine lessons were in flight at once, each failure was expensive and hard to read. A course is not one feature; it is this feature nine times plus a queue, an ordering, and an approval step per lesson.

So the unit is one video. It runs in about forty minutes, costs roughly half a dollar, and when it fails it fails once, legibly. A course is then a list of these, built one at a time, each approved before the next — which is how the content actually gets reviewed in practice anyway.

The seven stages

These are not new. They are the same stages our hand-made videos have always gone through — the same code that runs when a video is ordered in Slack. The page is a new door onto an existing building.

01 Research Turns the typed line into a brief: the outcome, the audience, the one scenario Ali is followed through. free · ~30s	02 Script Writes ~20 beats. One beat is one spoken sentence plus the picture shown while it is spoken. free · ~2min	03 Gate Reviews the script against the house standard and sends it back with line-level fixes until it passes. free · loops	
04 Produce Generates the illustrations, the voice, the motion, and renders the film with brand bumpers. PAID · ~25min	05 QA Scores it on seven factors. Below our threshold and it does not proceed. free · ~1min	06 Review Stops. A person watches the finished video and decides. HUMAN	07 Upload Publishes to YouTube, unlisted, and registers it for the LMS to pull. on approval

Two decisions, not one

Writing is free

Stages 1–3 are model calls on a credential we already pay for. Nothing is bought. So a person can type a topic, read the script, and throw it away at no cost — which means they will actually try things.

Making it costs money

Stage 4 buys generated art and generated speech. Measured at \$0.52 for a finished 20-beat video, against a \$1.77 estimate. It does not start until someone presses a second button, having read the script.

Nothing is spent by accident, and nothing is published by accident.

Two separate presses, with something real to read before each. The first shows you a script; the second shows you a video. Neither happens on its own.

Six things every script must do

RULE	WHAT IT MEANS IN THE SCRIPT
One concept	One thing a learner can do afterwards, written as an observable action. Not "understand"; nobody can check that.
One protagonist	Ali, always Ali, in one running situation followed all the way through — not a list of examples. A shopkeeper, a teacher, a clinic receptionist, whoever fits.
The five-part arc	Friction, then the fix, then the structure behind it, then how it breaks, then the payoff. In that order, on the same task.
Real numbers	"A five hundred rupee refund", not "a financial transaction". At least three data cards carrying actual figures.
One checkpoint	A question the learner answers — see overleaf. Exactly one, two-thirds of the way through.
Nothing written on screen	Illustrations carry no text, letters or numbers. Text is drawn by the renderer, where it can be spelled correctly and translated.

The gate sends work back; it does not just refuse it

A reviewer that can only say no is of limited use. When the gate finds a problem it returns line-level fixes naming the beat, and the writer produces a patch — only the beats that change, applied in code, so a redraft cannot quietly damage a beat nobody complained about.

The same is true of the automated checks inside stage 4. If the art judge rejects an illustration twice in a row, the defect is in the wording rather than the dice, so it asks for the beat to be rewritten instead of ending the run. Each reviewer gets its own budget of rounds.

"One blocker: beat 19 is a scene whose art prompt describes no physical act — it states both hands rest open and empty, which is a portrait."

A real gate finding. Specific enough to fix in one edit, which is the standard every reviewer in this system is held to.

What the reviewer at stage 6 is actually for

Everything before stage 6 is a machine checking a machine. Those checks are good — the QA score on our last video was 6.48 out of 7 — but they cannot tell whether a video is one Taleemabad would want to put its name on. That judgement is not automatable and we have not tried to automate it.

So the pipeline stops, hands over a finished video, and waits. Approve and it publishes. Do nothing and nothing happens.

The question is not in the video

Every script carries one multiple-choice question that is **never drawn and never spoken**. It occupies no time in the film. Its position in the script is a pause point: the player stops there, the LMS shows the question as a popup, the learner answers, and the video resumes.

This replaced an earlier design where the question was drawn on a card. That version taught the answer to a viewer who had no way to answer, and the platform learned nothing about whether anyone got it right. Moving it into the player makes it a real question — the learner answers, the LMS marks it, and the explanation is written for someone who has just chosen wrong.

The pause always lands between sentences

Each beat is one spoken sentence followed by a short silence, so a checkpoint sitting between two beats can never cut into a word. The API gives the exact second, measured from the start of the delivered file — including the brand intro, which is a common source of a two-second error.

What a learner sees

- The video plays normally and pauses on its own at the checkpoint.
- A panel appears: *"Let's answer a quick question, to check the concept landed."*
- The question, and four options. Three of them are mistakes a real learner makes.
- Right answer — a brief confirmation, and the video continues.
- Wrong answer — the explanation, which says why **their** choice was wrong and what the right one was. Then they continue.

The LMS is the system of record for attempts.

We serve the question; we do not store who answered what. Our service has no durable database, and scores belong with the learner's record, not with the content pipeline. The attempts endpoint deliberately returns "not implemented" rather than pretending to save.

What to expect, in numbers

	MEASURED	NOTE
Script written & passed	3–7 min	Longer when the gate asks for redrafts, which is normal and visible on the page.
Video produced	~25 min	Art, voice, motion, render, bumpers, then seven QA factors.
Cost per video	\$0.52	Generated art and speech. Writing and reviewing add nothing.
Finished length	1.5–2.5 min	20 beats at roughly six seconds of speech each.
QA score	6.48 / 7	Our publication threshold is 4.9. Seven factors, all rated.
Checkpoints per video	1	Exactly one, and it is mandatory.

One limitation to design around

Our service keeps no disk between deployments. A finished video survives until the next release, and its YouTube link survives permanently. **Treat the YouTube link and the checkpoint payload as the durable artefacts**, and store both on your side when a video is published. Do not rely on being able to fetch a video file back from us later.